

Lesson #7

Increasing system security with SELinux

Advanced Linux Administration

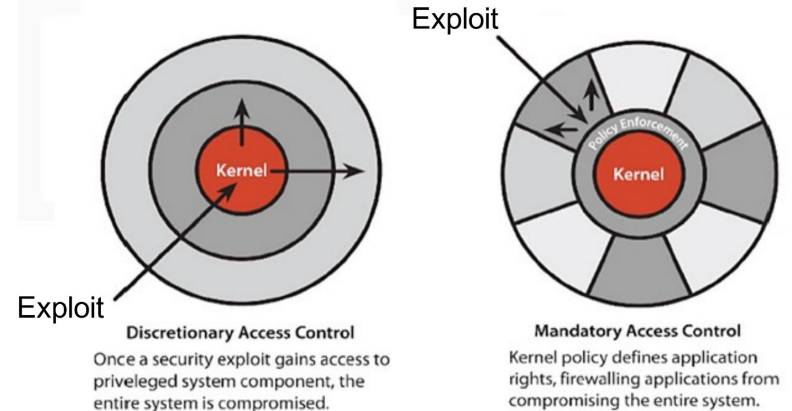
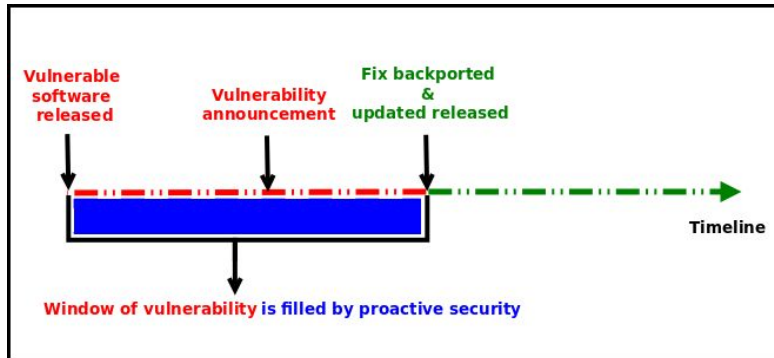
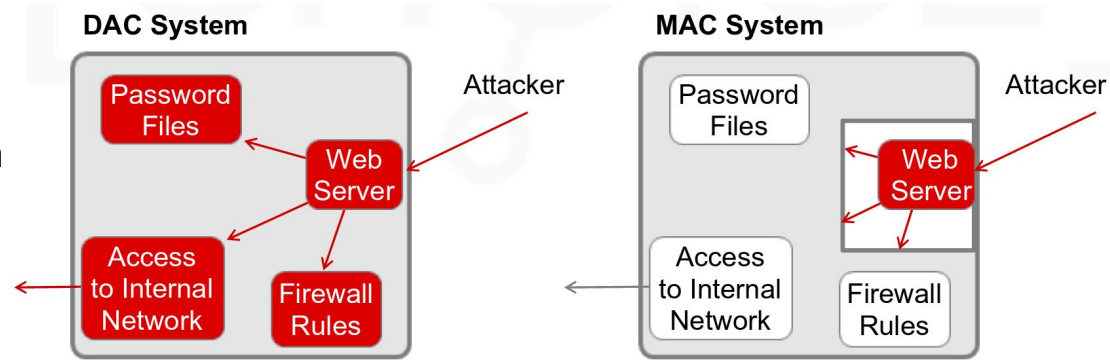
Karel Šrot, Eduard Beneš

Introduction

30 min



Motivation

- Compromised Service or Kernel
- Window of vulnerability
 - Reactive VS proactive approach
- Few exploits SELinux can prevent
 - Venom
 - Docker CVE-2016-9962
 - [Shellshock](#)



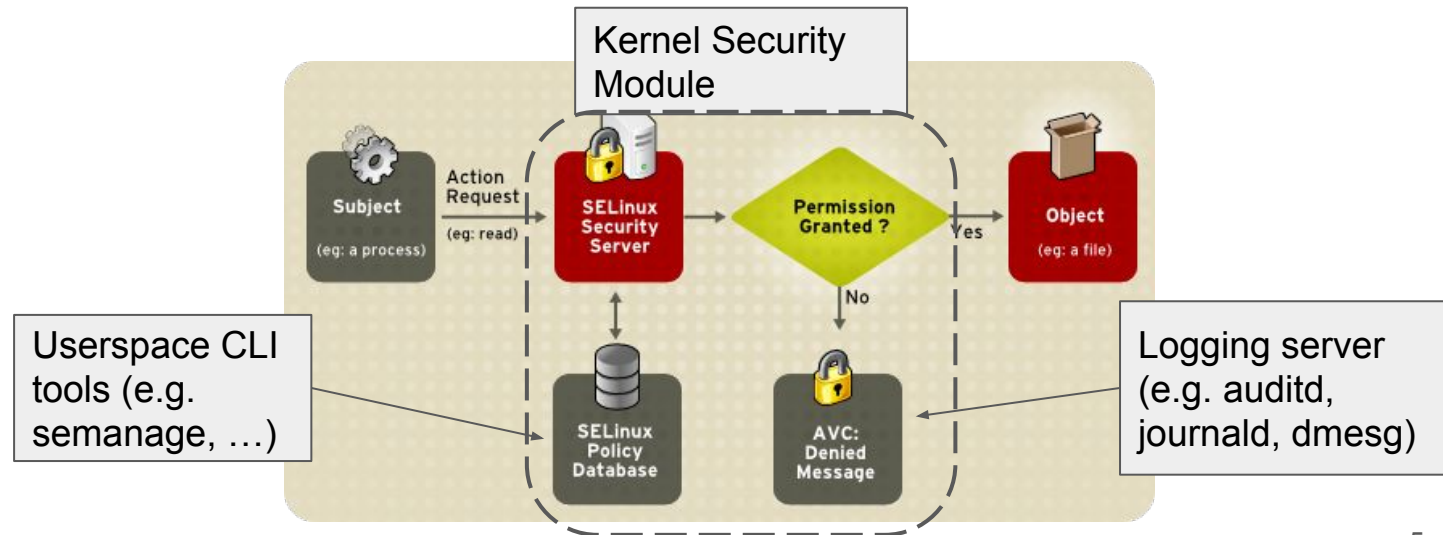
What is SELinux and how it can help?

- [SELinux](#) brings proactive security to your system with Mandatory Access Control (MAC).
- Provides isolation by separation of processes and resources on your system.
 - With process isolation it can mitigate attacks via [privilege escalation](#).
 - All you need to know about SELinux is in [this coloring book](#) :)
- It is enforced by a Kernel!
- All processes and files have a SELinux context (labels)
- SELinux policy is enforced system-wide
- Other known implementations of MAC are:
 - [AppArmor](#)
 - [Smack](#)



SELinux Architecture Overview

- Kernel security module and user-space tools.
- Security policy rules are loaded into policy database in kernel from user-space.
- Kernel enforces the security policy on the system between Subjects and Objects
- When access is denied it is called AVC denial and message is generated.



SELinux Decision Making in Detail

SELinux Modes

- Enforcing (deny by default)
- Permissive (evaluate and allow)

Default mode is set in /etc/selinux/config

DAC checks

- User-Group-Other
- Read-Write-Execute

```
$ ls -la
```

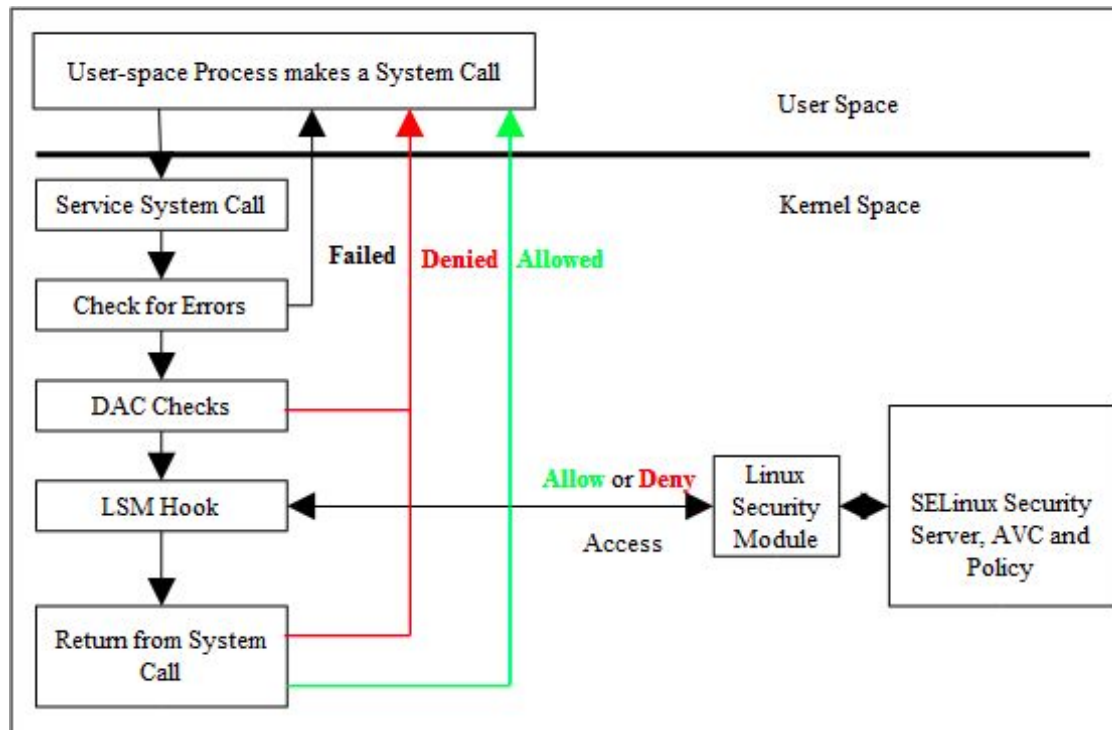
```
$ man chown
```

ACL

- Posix access control lists
- Allow fine-grained DAC rights on files and directories

```
$ getfacl /etc/passwd
```

```
$ man 5 acl
```



SELinux Labels (a.k.a. SELinux Context)

- SELinux context for all files (everything is a file) and processes is represented by a SELinux label
- Label consists of four fields:

<user>:<role>:<type>:<MLS/MCS>

- All of these information is used to make access control decisions by the Linux kernel
- Let's look at the labels on a system for different resources (file, process, port, ...)

```
$ ls -Z /etc/passwd
```

```
system_u:object_r:passwd_file_t:s0 /etc/passwd
```

```
# pstree -Z | head
```

```
# ls -Z /
```

```
# id -Z
```

```
# netstat -Z ; ss -Z
```

Type Enforcement (TE)

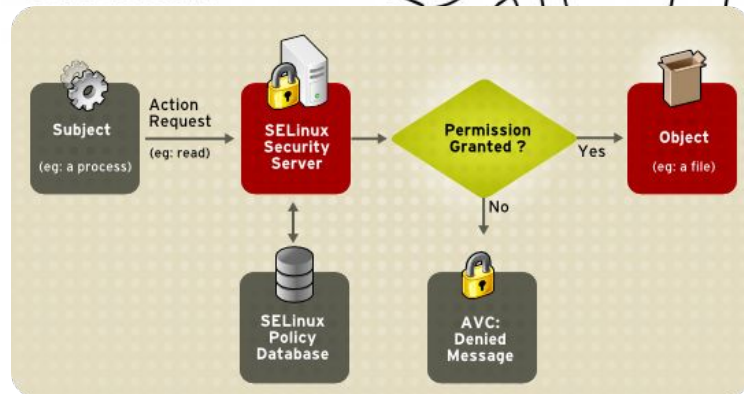

- Type Enforcement policy defines type of an subject and which action it can perform against object
- It uses just types (3rd field of a SELinux label) to define policy rules
- So called “unconfined” processes run with a types (e.g. unconfined_t) that are allowed almost all access, therefore they are not protected by SELinux when they get compromised.

`<user>:<role>:<type>:<MLS/MCS>`

```
# sesearch -s passwd_t -t etc_t -p write -A -ds -dt
```

TYPE ENFORCEMENT

Fido (dog:random1) trying to eat cat_chow:food is denied by type enforcement.

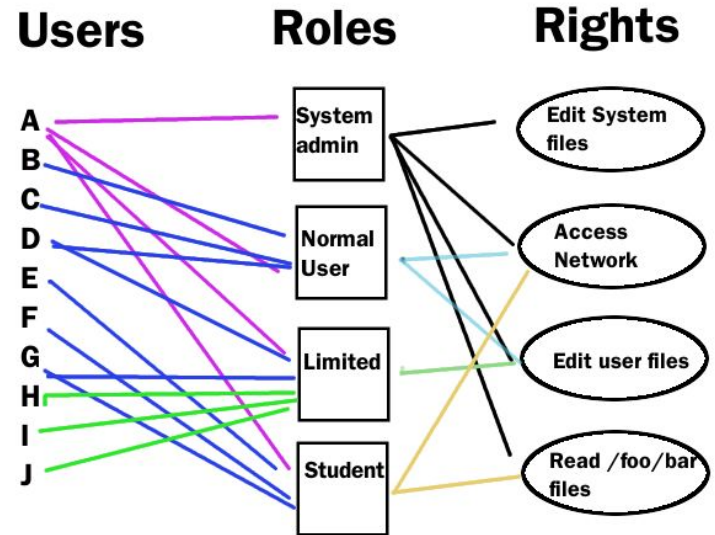


Role Based Access Control (RBAC)

- SELinux supports RBAC to implement refinement of security policy based on users and roles.
- The SELinux user definitions declare what the supported roles are that a user can "go" to
- Roles dictate which domains that the user (or session) is allowed to go in

<user>:<role>:<type>:<MLS/MCS>

```
# semanage user -l
# semanage login -l
# seinfo --role --user      (... -x)
# id -Z
# useradd -Z ...
# newrole -r sysadm_r
```



SELinux Policy Rules

- Rules are describing interactions between Subjects (process) and Objects (file, process, socket, ...)
- Rules can be also other than allow (dontaudit, neverallow)
- For example, we might have a rule like this to allow httpd_process read its log files
 - allow httpd_process httpd_log: FILE READ;
- You can easily generate custom policy rules using audit2allow from logged AVC denials
- General pattern of rules used to define SELinux policy loaded into kernel:

```
rule subject-type object-type : object-class { permissions }
```

- Real example following the rule pattern above:

```
# sesearch -s passwd_t -t etc_t -p write --allow -ds -dt  
allow passwd_t etc_t:file { append getattr ioctl read write };
```

SELinux Booleans

- Booleans enable runtime customization of SELinux policy to adjust it for a desired system configuration. Changes can be temporary or persistent.
- Conditional rules can be easily enabled or disabled by admin. For example to allow httpd daemon to serve web pages from user home directories.
- Checking status of available booleans?

```
$ getsebool -a
```

```
$ semanage boolean -l
```

```
$ man booleans
```

```
# dnf install selinux-policy-doc
```

```
$ apropos _selinux
```

```
# setsebool -P httpd_enable_homedir
```

Audit Daemon

- Auditd is an userspace daemon part of [Linux Auditing System](#)
 - Make sure it is running # `systemctl status auditd.service`
- By default all AVC messages end up in `/var/log/audit/audit.log` managed by Auditd
- In case auditd is not running, AVC will go to systemd journal (or `/var/log/messages`), or kernel ring buffer (`dmesg`)
 - with `systemd-journald-audit.socket` enabled AVCs always end up in journal
 - auditd logs should be queried using `ausearch` tool
- Tools like `sealert`, `audit2allow`, `audit2why` can help you diagnose SELinux denials
- Audit record types: AVC, USER_AVC, SELINUX_ERR, USER_AUTH, ...

```
# aureport
```

```
# man ausearch
```

```
# passwd --help >& /root/output.txt
```

```
# ausearch -m avc -ts recent
```

```
# systemctl kill auditd.service
```

```
# journalctl -S-300
```

AVC Messages and Searching SELinux Rules

- Access Vector Cache (AVC) messages are generated whenever access to an operation is granted or denied by kernel
 - By default only AVC denials are logged
 - To log allowed or hide denied operation `auditallow` and `dontaudit` rules can be used
- You can use `sesearch` from `setools-console` package to query loaded policy

```
# passwd --help >& /root/output.txt
# ausearch -c "passwd" ; ausearch -m avc -ts recent -sv no
----
time->Wed Oct 16 12:40:35 2019
type=AVC msg=audit(1571222435.365:8278): avc:  denied  { write } for  pid=8924
comm="passwd" path="/root/output.txt" dev="dm-1" ino=528771
scontext=unconfined_u:unconfined_r:passwd_t:s0-s0:c0.c1023
tcontext=unconfined_u:object_r:admin_home_t:s0 tclass=file permissive=0

# sesearch -p write -s passwd_t -t admin_home_t --allow
```

Checking SELinux Status on a System

- RTFM .. man pages
- Is SELinux in Enforcing or Permissive mode? Does it use targeted policy or some other?
getenforce
sestatus
- Switching to enforcing mode
setenforce 1
- In case of a mislabeled filesystem you can force to relabel during a next boot.
touch /.autorelabel; reboot
.. NOTE: this might take a looong time
- What is a default SELinux policy state?
cat /etc/selinux/config
- I need more information
seinfo
avcstat

Disabling SELinux

- No longer possible through /etc/selinux/config
- Update kernel parameters with selinux=0

```
# grubby --update-kernel  
ALL --args selinux=0
```

- To enabled SELinux again

```
# grubby --update-kernel  
ALL --remove-args selinux
```

Managing File SELinux Context

- SELinux policy defines context for newly created objects depending on its type (-m) and path
- Context is by default 1) inherited or 2) policy default is used when creating new objects

```
$ matchpathcon -m type /some/path/to/dir/or/file
```

- **Temporary** changes with chcon will not survive reboot or restoring to default context

```
# mkdir /web/  
# chcon -R -t httpd_sys_content_t /web/  
# restorecon -R -v /web/
```

- **Persistent** changes to the policy will survive also reboot. Using `semanage fcontext` changes will be saved to /etc/selinux/targeted/contexts/files/ directory, (see man pages)

```
# semanage fcontext -a -t httpd_sys_content_t "/web(/.*)"?"  
# restorecon -R -v /web
```

- **Restore** file context to its default according to current SELinux policy:

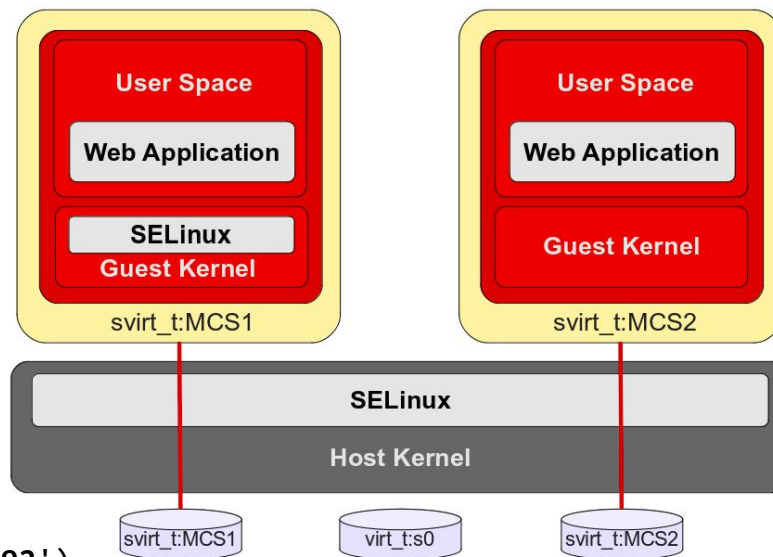
```
# restorecon -v some-file ... (use -n for passive check without applying any changes)  
# restorecon -R -v directory-name  
# man fixfiles
```

Usage of MCS for Virtual Guests and Containers

- Type enforcement on host system using just dedicated types `container_t`, `container_runtime_t`, `container_file_t` is not sufficient to isolate running container processes
- Unique MCS category is used to confine containers and protect them from each other. Kernel enforces that MCS must match after Type is enforced
- Each container process is assigned an unique MCS

```
# pstree -Zp | grep container_t
|   |-bash(24031,`system_u:system_r:container_t:s0:c137,c292')
|   |-bash(2597,`system_u:system_r:container_t:s0:c563,c1010')
|   |-bash(2681,`system_u:system_r:container_t:s0:c119,c927')
```

```
# ls -lZ `podman inspect -l --format "{.HostnamePath}"`
-rw-r--r--. 1 ... system_u:object_r:container_file_t:s0:c119,c927 12 Oct 15 09:47
/run/user/1000/overlay-containers/64c02f69cac16bc360b06a4ff8d2181a57b5a273a89412bae0533ac1e
dbf3c6a/userdata/hostname
```



Overview of SELinux policy management tools

- Use man pages for detailed description, options, and usage examples
- File contexts
 - `# semanage fcontext`
- Booleans
 - `# semanage boolean`
- Ports
 - `# semanage port`
- Permissive domains
 - `# semanage permissive`
- Searching policy and logs
 - `# sesearch`
 - `# ausearch`
- Creating local policy modules by parsing policy logs
 - `# audit2allow`
- Building and managing policy modules
 - `# semodule`

Workshop

60 min

Workshop labs

- In the following next slides there are 7 labs total
 - 6 regular + 1 bonus lab
- Each lab has a time estimate how much time should you spend on it if everything goes well
- We encourage you to help each other or raise your hand to get help from lecturers
- If you couldn't finish all labs during this class, you should complete them on your own later because learned skills will be used in following lectures or during a final practical exam.
- *HINT: Focus on the lab content and leave any exploration or deep dive desires as a self-study for later*
- *HINT2: Use or create a fresh snapshot of your virtual guest machine to have and preserve a clean work environment without breaking your system during the labs.*

Overview of SELinux Labs

1. Lab1 - SELinux status and AVC logs (5min)
2. Lab2 - Basic policy management (5 mins)
3. Lab3 - Searching audit logs (5min)
4. Lab4 - Changing file context temporarily (5min)
5. Lab5 - Permanent policy changes (10m)
6. Lab6 - Simple SELinux booleans (10min)
7. Bonus Lab - Allow httpd access to homedirs (15min)

Lab1 - SELinux status and AVC logs ½ (5min)

- Check status of SELinux on your system

`getenforce ; sestatus` → check it is in Enforcing mode using Targeted policy

`cat /etc/selinux/config`

- Install packages required for all labs

`dnf install selinux-policy-doc selinux-policy-devel policycoreutils-python
policycoreutils-python-utils setools-console`

Still you may have to install some additional packages later.

Lab1 - SELinux status and AVC logs 2/2 (5min)

- Finding logged SELinux AVC denials with(&out) auditd service running

systemctl status auditd.service → make sure it is active and enabled

passwd --help >& /root/output.txt → this will generate AVC denial

cat /root/output.txt ; → this should be empty

ausearch -m avc -sv no -ts recent → there should be AVC denial message for the passwd

- Now kill the auditd service and do the same action with the passwd

systemctl kill auditd.service → make sure it is not running with systemctl status

passwd --help >& /root/output.txt

cat /root/output.txt ; ausearch -m avc -sv no -ts recent

- All SELinux logs/denials should now end up in systemd journal

journalctl -S-60 → ... please raise your hand if not! :)

Lab2 - Basic policy management 1/2 (5 mins)

- Set the SELinux to Permissive mode using the Targeted policy

`setenforce permissive` → check it is Permissive mode with commands from previous lab

- We want our logs to go to `audit.log` so better reboot

`systemctl status auditd.service` → make sure it is active and enabled

- Save system time for checking logs and execute previously denied `passwd` action again.

`d1=`date +%T``

`passwd --help >& /root/output.txt`

`cat /root/output.txt`

`ausearch -m avc -sv no -ts $d1` → make sure you've defined the `d1` variable

- Now, compare logged audit records to the output from Lab1 when in Enforcing mode
- There should be a new record for { `ioctl` } together with the previous { `write` }, and you should see `permissive=1` at the end of the audit record compared to `permissive=0` before (look at the outputs from lab1)

Lab2 - Basic policy management 2/2 (5 mins)

- The passwd action in Permissive mode should be allowed, all generated AVC denials recorded in logs and it should print the help message. Please let us know if something went wrong! :)
- Now set the system back to Enforcing mode and make sure auditd.service is running

`systemctl status auditd.service` → start the service if it is not running properly

`setenforce 1`

`getenforce`

Lab3 - Searching audit logs ½ (5min)

Using ausearch tool we'll be working with previously generated events in `/var/log/audit/audit.log` and practicing various ways to search the logs.

```
# man ausearch
```

- Search the audit logs by name of executable using `-c` option and try `--pid` to specify process

```
# ausearch -c passwd → there should be all the denials you've generated in lab1,2
```

... now try to create a search using the `--pid` option

- Next you can specify a time period for the search using `-ts time-start -te time-end`, for example:

```
# ausearch -ts 13:30 -te 13:55 → ... please adjust the time values to for your class hours ;-)
```

- Or search by a file name:

```
# ausearch --file /root/output.txt
```

- Search based on success value. SELinux AVC denials have success value **no**

```
# ausearch -m avc -sv no
```

Lab3 - Searching audit logs 2/2 (5min)

- Audit records can have various types, .e.g. AVC, USER_AUTH, USER_LOGIN, ... let's try them! :)

```
# ausearch -m user_auth -sv no
```

- Sometimes you might need to search audit logs related to a specific SELinux context

```
# ausearch -se container_t
```

- Now, authenticate successfully using ``$ sudo -i`` as a normal user and then as a user root list the corresponding success audit record:

```
user$ sudo -K ; sudo -i
```

```
root# ausearch -m user_auth -sv yes -ts recent -x /usr/bin/sudo
```

```
----
```

```
time->Fri Oct 18 09:31:05 2019
```

```
type=USER_AUTH msg=audit(1571383865.850:8689): pid=4882 uid=1000 auid=1000 ses=3
```

```
subj=unconfined_u:unconfined_r:unconfined_t:s0-s0:c0.c1023 msg='op=PAM:authentication grantors=pam_unix acct="ebenes"
```

```
exe="/usr/bin/sudo" hostname=? addr=? terminal=/dev/pts/9 res=success'
```

Lab4 - Changing file context temporarily 1/2 (5min)

In this lab we'll adjust context of `/root/output.txt` to allow access to `passwd`.

```
# systemctl status auditd.service ; getenforce → make sure we are good to go
```

- Compare context of these two files and set context of the `output.txt` to match `/etc/passwd`

```
# ls -lZ /etc/passwd /root/output.txt
```

```
# chcon -v --reference=/etc/passwd /root/output.txt
```

```
# restorecon -v /root/output.txt → to set it back to default context
```

```
# chcon -v -t passwd_file_t /root/output.txt → try to change it again
```

```
# ls -lZ /etc/passwd /root/output.txt → they should have the same context
```

- Now lets try to run `passwd` example again and observe it to succeed! :)

```
# d1=`date +%T` ; passwd --help >& /root/output.txt
```

- There should be no AVC denials related to the action before.

```
# cat /root/output.txt
```

```
# ausearch -m avc -sv no -ts $d1
```

Lab4 - Changing file context temporarily 2/2 (5min)

- Finally put the system into proper state. First our file and then recursively on a directory.

```
# restorecon -v /root/output.txt
```

```
Relabeled /root/output.txt from system_u:object_r:passwd_file_t:s0 to system_u:object_r:admin_home_t:s0
```

```
# restorecon -R -v /root
```

```
# rm /root/output.txt
```

Lab5 - Permanent policy changes 1/2 (10m)

In this lab we'll change context of `/root/output.txt` but permanently so it will survive reboot and relabelling of a system. Similarly you could change a label of port, boolean, ...

```
# man semanage-fcontext
```

- Change system policy to label `/root/output.txt` as `passwd_file_t` and let it to fix the label.

```
# semanage fcontext -a -t passwd_file_t /root/output.txt
```

```
# restorecon -v /root/output.txt → you should be notified about the relabel
```

```
# ls -lZ /etc/passwd /root/output.txt → now, they should have the same context
```

- Now let's try to run `passwd` example again and observe it to succeed! :)

```
# date ; passwd --help >& /root/output.txt
```

- There should be no AVC denials logged as related to the `passwd` after time returned by `date`

```
# cat /root/output.txt
```

```
# ausearch -m avc -sv no -ts recent
```

- ~~• At home you can try relabel the whole system during reboot ... take a break :)~~

```
# touch /.autorelabel ; reboot
```

Lab5 - Permanent policy changes 2/2 (10m)

- Check that the SELinux type of `output.txt` still matches the context of `/etc/passwd`

```
# ls -lZ /etc/passwd /root/output.txt
```

- Finally remove the persistent rule and cleanup the system

```
# semanage fcontext -d -t passwd_file_t /root/output.txt
```

```
# restorecon -v /root/output.txt → you should be notified about the relabel
```

```
Relabeled /root/output.txt from system_u:object_r:passwd_file_t:s0 to system_u:object_r:admin_home_t:s0
```

```
# rm /root/output.txt
```

Lab6 - Simple SELinux Booleans 1/2 (5min)

In this lab we'll demonstrate usage and effect of SELinux booleans to a ping command for user_u.

```
# systemctl status auditd.service ; getenforce → make sure we are good to go
```

- List all available booleans on your system

```
# getsebool -a | less
```

```
# semanage boolean -l | less
```

- First we'll need to create a user with user_u role and login via ssh to use it.

```
# useradd -Z user_u useru
```

```
# passwd useru
```

```
# semanage login -l
```

- ... you might need to install and enable sshd.service first

```
# ssh useru@localhost
```

```
useru$ id -Z → are you really ssh-ed as a user with SELinux user_u ?
```

```
useru$ ping localhost → this should work without any issues
```

Lab6 - Simple SELinux Booleans 2/2 (5min)

```
useru$ ping localhost
```

```
...
```

- Now find the related selinux boolean and turn it off.

```
# semanage boolean -l | grep ping
```

```
# setsebool -P selinuxuser_ping off
```

- Login again as the confined user and try to use the ping or traceroute commands again

```
useru$ ping localhost → should give you permission denied ... check for AVC denials
```

```
# sesearch -b selinuxuser_ping --allow → this is list of allow rules used by this boolean
```

- Finally list all local boolean customizations and return the ping boolean value to an default state

```
# semanage boolean -l -C
```

```
# setsebool -P selinuxuser_ping on
```


Extra Lab7 - Allow httpd access to homedirs (15-20 min)

In this lab we'll be using SELinux booleans to allow httpd server to access user home directories.

- First try to find httpd related SELinux booleans and man pages

```
# getsebool -a | grep http
```

```
# semanage boolean -l | grep http
```

```
# apropos http → .. you might need to run the `mandb -c` first
```

- Follow instructions in the following HOWTO ;-)

<https://www.if-not-true-then-false.com/2010/enable-apache-userdir-with-selinux-on-fedora-centos-red-hat-rhel/>

- **WITH THE FOLLOWING CORRECTIONS:**

- Step 6) instead of `chcon -R ... do`

```
# restorecon -Rv /home/testuser/public_html
```

- After that when creating `test.html` and `test.php` files do

```
# chown -R testuser:testuser /home/testuser/public_html
```

Links and References

- [short comparison and reference to other authorization models](#)
- Shellshock exploit Demo - <https://www.youtube.com/watch?v=Ysshrh4aGOs#action=share>
- SELinux [coloring book](#)
- <https://wiki.centos.org/TipsAndTricks/SelinuxBooleans>
- https://access.redhat.com/documentation/en-us/red_hat_enterprise_linux/8/html-single/using_selinux/index
- https://access.redhat.com/documentation/en-us/red_hat_enterprise_linux/7/html/selinux_users_and_administrators_guide/
- https://access.redhat.com/documentation/en-us/red_hat_enterprise_linux/8/html-single/using_selinux/index#customizing-the-selinux-policy-for-the-apache-http-server-in-a-non-standard-configuration_configuring-selinux-for-applications-and-services-with-non-standard-configurations