

Lesson #10 - systemd unit file management

Advanced Linux Administration

Aleš Zelinka

Introduction

30 min

Goals

- Creation of a new unit file for a service
- Customization of an existing unit file
- Grouping of services into targets
- Managing resources used up by units (cgroups)
- Non-service unit types

Quick recap of systemd basics from lesson #2

- Everything is a unit: service / target / ...
- `systemctl status|start|stop|enable|disable`
- Units have dependencies on each other and through this create a tree
 - `$ systemctl list-dependencies <unit>`

Unit vs Unit file

- systemd unit is an abstract object, a basic building block systemd works with
- unit file is a config file, a textual representation of the unit

- sshd.service ~ the unit
 - `$ systemctl status unit`
- `/usr/lib/systemd/system/sshd.service`
 - `$ man systemd.unit`
 - `$ man systemd.syntax`

```
zelial@chuchvalec:~  
File Edit View Search Terminal Help  
[chuch ~]$ sudo systemctl status cups  
● cups.service - CUPS Scheduler  
   Loaded: loaded (/usr/lib/systemd/system/cups.service; enabled; vendor preset: disabled)  
   Active: active (running) since Wed 2019-09-18 19:43:12 CEST; 40min ago  
     Docs: man:cupsd(8)  
  Main PID: 1070 (cupsd)  
    Status: "Scheduler is running..."  
      Tasks: 1 (limit: 4915)  
   CGroup: /system.slice/cups.service  
           └─1070 /usr/sbin/cupsd -l  
  
Sep 18 19:43:12 chuchvalec cupsd[1070]: Loading from cache...  
Sep 18 19:43:12 chuchvalec systemd[1]: Started CUPS Scheduler.  
Sep 18 19:43:12 chuchvalec cupsd[1070]: REQUEST localhost - - "POST / HTTP/1.1" 200 349 Create-Printer-Subscription  
Sep 18 19:43:13 chuchvalec cupsd[1070]: REQUEST localhost - - "POST / HTTP/1.1" 200 176 Create-Printer-Subscription  
Sep 18 19:43:26 chuchvalec cupsd[1070]: REQUEST localhost - - "POST / HTTP/1.1" 200 362 Create-Printer-Subscription  
Sep 18 20:22:37 chuchvalec cupsd[1070]: REQUEST localhost - - "POST / HTTP/1.1" 401 123 Cancel-Subscription success  
Sep 18 20:22:37 chuchvalec cupsd[1070]: REQUEST localhost - root "POST / HTTP/1.1" 200 123 Cancel-Subscription succ  
Sep 18 20:22:37 chuchvalec cupsd[1070]: REQUEST localhost - - "POST / HTTP/1.1" 200 152 Cancel-Subscription success  
Sep 18 20:22:37 chuchvalec cupsd[1070]: REQUEST localhost - - "POST / HTTP/1.1" 200 349 Create-Printer-Subscription  
Sep 18 20:22:37 chuchvalec cupsd[1070]: REQUEST localhost - - "POST / HTTP/1.1" 200 176 Create-Printer-Subscription  
lines 1-20/20 (END)
```

Anatomy of a service unit file

A unit file is a plain text ini-style file

- **[Unit]** section - common unit options used during runtime
 - `$ man systemd.unit`
- **[Service]** section - specific to services
 - `$ man systemd.service`
- **[Install]** section - common unit options used when enabling/disabling the unit (during 'installation')
 - Tells systemd where to put the unit in the dependency tree
 - `$ man systemd.unit`

```
[Unit]
Description=OpenSSH server daemon
Documentation=man:sshd(8) man:sshd_config(5)
After=network.target sshd-keygen.target
Wants=sshd-keygen.target

[Service]
Type=notify
EnvironmentFile=/etc/crypto-policies/back-ends/opensshserver.config
EnvironmentFile=/etc/sysconfig/ssh
ExecStart=/usr/sbin/sshd -D $OPTIONS $CRYPTO_POLICY
ExecReload=/bin/kill -HUP $MAINPID
KillMode=process
Restart=on-failure
RestartSec=42s

[Install]
WantedBy=multi-user.target
```

Anatomy of a service unit file - [Service] section #1

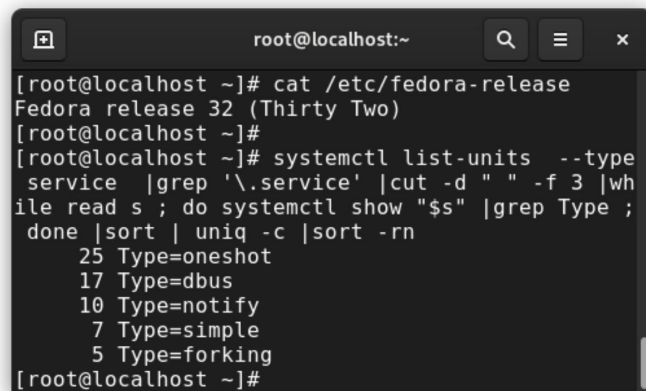
Q: How does systemd know that a service is started / running?

Anatomy of a service unit file - [Service] section #2

Important options: **Type**

How does systemd know that a service is started?

- **Simple** - default. Considered started after the process is forked. For simple, single-process stuff.
- **Notify** - the service tells systemd
- **Forking** - classic unix daemons, considered started after process has exited (after forking a child process)
- **Oneshot** - considered started after the process exits. For stuff that doesn't keep running
- ...

A terminal window titled 'root@localhost:~' with search, menu, and close buttons. It shows the output of several commands: 'cat /etc/fedora-release' returns 'Fedora release 32 (Thirty Two)'. 'systemctl list-units --type service | grep '\.service' | cut -d " " -f 3 | while read s ; do systemctl show "\$s" | grep Type ; done | sort | uniq -c | sort -rn' returns a list of unit types and their counts: 25 Type=oneshot, 17 Type=dbus, 10 Type=notify, 7 Type=simple, and 5 Type=forking.

```
root@localhost:~  
[root@localhost ~]# cat /etc/fedora-release  
Fedora release 32 (Thirty Two)  
[root@localhost ~]#  
[root@localhost ~]# systemctl list-units --type  
service | grep '\.service' | cut -d " " -f 3 | wh  
ile read s ; do systemctl show "$s" | grep Type ;  
done | sort | uniq -c | sort -rn  
    25 Type=oneshot  
    17 Type=dbus  
    10 Type=notify  
     7 Type=simple  
     5 Type=forking  
[root@localhost ~]#
```


Anatomy of a service unit file - [Service] section #3

Important options: **ExecStart** & friends

- **ExecStart** - what gets executed when systemctl start is called
 - `ExecStart=/usr/bin/my_service`
- **ExecStop** - command responsible for stopping the service, if not specified service gets killed (SIGTERM + SIGKILL)
- **ExecReload** - command used to reload configuration, if not specified `systemctl reload <service>` will not be available

Stdout & stderr goes to journal ~ `$ journalctl -u name.service`

These are not shell commands! You can't use shell constructs directly



```
ExecStart=test && echo yes || echo no
```



```
ExecStart=bash -c "test && echo yes || echo no"
```

Anatomy of a service unit file - [Service] section #4

For passing options to the executed process without having to edit the unit file itself.

Q: Why should I not directly edit the unit file?

- **Environment** - define environmental variables directly
 - `Environment=LANG=C "OTHER_VAR=a b c d"`
- **EnvironmentFile** - path to a file where environment variables are defined
 - `EnvironmentFile=/etc/my_service_config`

```
/usr/lib/systemd/system/sshd.service
...
[Service]
EnvironmentFile=/etc/crypto-policies/back-ends/opensshserver.config
EnvironmentFile=/etc/sysconfig/ssh
ExecStart=/usr/sbin/sshd -D $OPTIONS $CRYPTO_POLICY
ExecReload=/bin/kill -HUP $MAINPID
...
```

Anatomy of a service unit file - Documentation

[Unit]

Description - A human readable name for the unit.

- `Description=OpenSSH server daemon`

Documentation - A space-separated list of URIs referencing documentation for this unit or its configuration

- `Documentation=man:sshd(8) man:sshd_config(5)`
- `Documentation=man:auditd(8) https://github.com/linux-audit/audit-documentation`

Anatomy of a service unit file - relations

There are two types of relations in between units

- Activation
 - A unit X has to be started before dependant unit Y can be started
 - E.g. a service wants to be autostarted after boot so it makes the default target depend on itself
- Ordering
 - When both X and Y units are started (e.g. because of their activation dependencies), X has to start before Y - Y's start is delayed after X has finished starting
 - E.g. a web server service only starts after network-online.target is reached

Anatomy of a service unit file - activation dependency #1

[Install] section

- **WantedBy**=foo.service in a service bar.service means that bar will be started when foo is being started.
 - Bar will be started even if foo fails to start
- **RequiredBy**=foo.service in a service bar.service means that bar will be started when foo is being started.
 - Bar will **not** be started if foo fails to start

```
[Install]  
WantedBy=multi-user.target
```

Anatomy of a service unit file - activation dependency #2

Wants and **Requires** are reversed **WantedBy** and **RequiredBy** used in the [Unit] section:

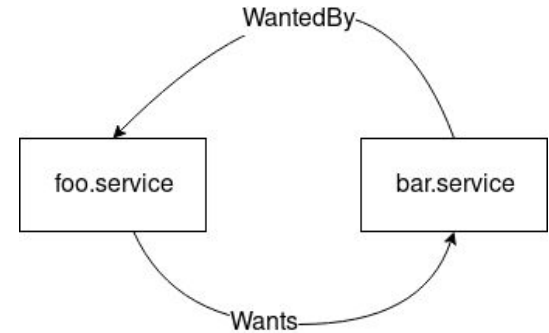
- `bar.service` having `Wants=foo.service`

is equivalent to

- `foo.service` having `WantedBy=bar.service`

`$ systemctl list-dependencies`

- Shows the tree based on activation dependencies



Anatomy of a service unit file - ordering #1

[Unit] section

- **After/Before** - enforce ordering of unit starting

```
[Unit]
Description=OpenSSH server daemon
Documentation=man:sshd(8) man:sshd_config(5)
After=network.target sshd-keygen.target
```

\$ systemctl --after list-dependencies

- Shows the tree based on ordering dependencies

Editing a unit file

Do not edit vendor unit files in `/usr/lib/systemd/system/!`

- New unit file or whole replacement @ `/etc/systemd/system/`
 - `$ systemctl edit --full name.unit`
- Drop-in snippet @ `/etc/systemd/system/name.unit.d/`
 - All snippets get merged with the vendor unit file
 - `$ systemctl edit name.unit`
- Drop all changes, reverting to the vendor file
 - `$ systemctl revert name.unit`
- Show what and from where will systemd load content
 - `$ systemctl status|cat name.unit`


DEMO: edit `httpd.service`, show status & cat

Managing resources of a service: cgroups

Systemd tracks all processes belonging to a service via [cgroups](#).

cgroups allows you to

- Hierarchically group and label processes
 - Systemd enforces flat hierarchy of a unit == cgroup
- Apply resource limits to the groups
 - Memory
 - CPU
 - IO
 - Network (ipv4 & ipv6)
 - Tasks (processes and threads)



Systemd + cgroups: tooling

```
$ ps axwf -eo pid,user,cgname:50,args
```

Systemd specific tooling for working with cgroups

- `$ systemctl status name.unit`
- `$ systemd-cgls`
- `$ systemd-cgtop`

Control Group	Tasks	%CPU	Memory	Input/s	Output/s
/	392	-	1.2G	-	-
/init.scope	1	-	68.1M	-	-
/system.slice	288	-	957.3M	-	-
/system.slice/ModemManager.service	3	-	8.2M	-	-
/system.slice/NetworkManager.service	4	-	19.0M	-	-
/system.slice/abrt-journal-core.service	1	-	10.9M	-	-
/system.slice/abrt-oops.service	1	-	8.2M	-	-
/system.slice/abrt-xorg.service	1	-	17.1M	-	-
/system.slice/abrttd.service	3	-	37.3M	-	-
/system.slice/alsa-state.service	1	-	352.0K	-	-
/system.slice/atd.service	1	-	416.0K	-	-
/system.slice/auditd.service	2	-	2.5M	-	-

Systemd + cgroups: configuration #1

You can turn a resource accounting from either a unit or globally.

Unit file, [Service] section	Global config (/etc/systemd/system.conf)
CPUAccounting=yes	DefaultCPUAccounting=yes
MemoryAccounting=yes	DefaultMemoryAccounting=yes
TasksAccounting=yes	DefaultTasksAccounting=yes
IOAccounting=yes	DefaultIOAccounting=yes
IPAccounting=yes	DefaultIPAccounting=yes

Reload systemd after any config change: `$ systemctl daemon-reload`

Systemd + cgroups: configuration #2

Each of CPU/Memory/Task/... has its own set of options that **limit** that particular resource consumption. See for details:

```
$ man systemd.resource-control
```

In [Service] section:

- **CPUQuota=20%** ~ service will use max 20% CPU
- **CPUWeight=[1..100..10000]** ~ proportional CPU usage, 100 is default => 200 means twice the CPU cycles
- **MemoryHigh=200M** / **MemoryHigh=20%** ~ limit before service will be preferred to be swapped out
- **MemoryMax=...** ~ limit before service will be forced swapped and if that's not enough, OOM killed
- **TasksMax=N** ~ maximum number of tasks (processes / threads) service can contain
- ...

Other unit types / activation modes

Unit files name.type that activates name.service when they are themselves activated.

- **Socket** ~ network connection activation. Defines a socket, once client connects to it, a corresponding service unit is started
- **Device** ~ HW activation. device units get created when HW is detected -> services can depend on them and get activated once HW is available
- **Mount** ~ mount a filesystem. Equivalent to a line from /etc/fstab
- **Path** ~ monitors a path (directory or file), triggers service on creation or change
- **Timer** ~ time-based activation. Activates after a specified interval relative to boot/unit activation or on a defined time/date
- `$ systemctl list-units --type=<type>`

Workshop

60 min

Workshop labs

- In the following next slides there are 5 labs total
 - 4 regular + 1 bonus
- Each lab has a time estimate how much time should you spend on it if everything goes well
- **We encourage you to help each other or raise your hand to get help from lecturers**
- If you couldn't finish all labs during this class, you should complete them on your own later because learned skills will be used in following lectures or during a final practical exam.
- *HINT: focus on the lab content and leave any exploration or deep dive desires as a self-study for later*

Lab 1 - Create a new service [10 min]

Create a service that prints current date/time every minute:

```
$ while : ; do date ; sleep 60; done
```

1. Create `/etc/systemd/system/my_service.service`
2. Use following options with correct values in appropriate sections ([Unit] vs [Service]):
 - a. Description
 - b. ExecStart
 - c. Type
3. Start the service
4. Verify that it's running
 - a. `$ systemctl status my_service`
5. Verify that it logs the time & date by observing its journal
 - a. `$ journalctl -f -u my_service`

Lab 2 - Make your new service start after boot [5 min]

Make the **my_service** service start after boot.

1. Stop the service if it's running
2. Edit the service using `$ systemctl edit ...`
3. Make the service start after boot by having it WantedBy multi-user.target. This belongs to the Install section.
4. Enable the service: `$ systemctl enable my_service`
 - a. This will make the 'WantedBy' option to take effect
5. Verify that the service has become part of the boot process (multi-user.target)
 - a. By listing everything that is part of the multi-user.target
 - i. `$ systemctl list-dependencies multi-user.target`
 - b. Alternatively by looking at everything that your service is part of
 - i. `$ systemctl list-dependencies --reverse my_service.service`
 - c. By activating the multi-user.target and checking that the service has been started
 - i. `$ systemctl start multi-user.target`

Lab 3 - Modify existing service [10 min]

Modify the **httpd** service so that it is automatically restarted if it crashes.

1. Start httpd service, check it's status: has to be running
2. Kill all its processes, simulating a crash `$ killall httpd`
3. Check status, should be inactive (dead)
4. Edit its unit file, use `$ systemctl edit httpd.service` to not alter the original file but create a drop-in override file instead
 - a. add `Restart=always` to the `[Service]` section
5. Restart the service
6. Verify with status that
 - a. It's running
 - b. The drop-in file is listed
7. Kill all the processes again
8. Verify with status that it is still running (and PIDs are different)
9. Revert the changes with `$ systemctl revert httpd.service`

Lab 3/1 - Enforce correct service order [15 min]

Create 3 separate services called **a**, **b** and **c**. Make sure they are always started in the alphabetical order.

1. Create 3 new minimal services called a, b and c using the template from this slide (change the string after echo to reflect the name of the service)
2. Modify a.service to introduce delay:
 - a. `ExecStart=bash -c "sleep 1; echo aaa"`
3. Enable all 3 services, making them part of the multi-user target
4. Test the services by activating the target they are part of
 - a. `$ systemctl start multi-user.target`
5. Verify they were started by checking their journal
 - a. `$ journalctl -f -u a -u b -u c | egrep '(echo)|(bash)'`
 - b. Preferably keep this running in a separate terminal
6. Verify in the journal that the 3 echos are being printed out of order (a is last)

```
[Unit]
[Service]
ExecStart=echo aaa
Type=oneshot
[Install]
WantedBy=multi-user.target
```

Lab 3/2 - Enforce correct service order [15 min]

Alter the services so that they always start in alphabetical order

1. Add the correct **After**=name.service (and/or **Before**=name.service) options to the [Unit] sections of the 3 services to enforce alphabetical order
2. Test by rerunning them
 - a. `$ systemctl start multi-user.target`
3. Verify correct ordering by looking at their logs
 - a. `$ journalctl -f -u a -u b -u c | egrep '(echo)|(bash)'`
4. Clean up: disable the 3 services and delete their unit files

Lab 4 - Limiting CPU usage [10 min]

Create a service that is very CPU intensive, limit it to maximum 10% CPU.

1. Create service **cpueater**
2. Start the service, observe that it uses all available CPU
 - a. openssl process in `$ top`
 - b. `cpueater.service` in `$ systemctl-cgls`
3. Edit the service, add to the [Service] sections options that will
 - a. Turn on CPU accounting
 - b. Set CPU quota to 10%
4. Restart the service
5. Verify the CPU usage again, should be 10% or lower
6. Stop the service, remove the unit file

```
/etc/systemd/system/cpueater.service

[Unit]
Description=Uses as much CPU as it can
[Service]
ExecStart=openssl speed
Restart=always
```

Bonus Lab 1 - Limiting Memory usage [?? min] 2 points

Similarly to previous lab: create a service that allocates at least X units of memory. Limit it using systemd to only be able to use X/2. E.g. let it allocate 1GB, limit it to 512MB.

- You need to come up with the tool/script/program/... that will consistently allocate set amount of memory
- You need to find out the correct systemd options that control memory usage.
 - Hint: check these slides or `$ man systemd.resource-control`